

**AAHAARA360 — A COMPREHENSIVE CLOUD-BASED
PRACTICE MANAGEMENT AND NUTRIENT ANALYSIS
PLATFORM FOR AYURVEDIC DIETITIANS ON CLOUD-
NATIVE ARCHITECTURE**

CLOUD COMPUTING

Submitted by

ATCHAYA S (230701045)

MADHUMITHA P(230701169)

MANISHAA G (230701176)

in partial fulfilment for the award of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING



RAJALAKSHMI ENGINEERING COLLEGE

ANNA UNIVERSITY, CHENNAI

MAY 2026

BONAFIDE CERTIFICATE

Certified that this project report titled "**AAHAARA360 — A COMPREHENSIVE CLOUD-BASED PRACTICE MANAGEMENT AND NUTRIENT ANALYSIS PLATFORM FOR AYURVEDIC DIETITIANS**" is the Bonafide work of **ATCHAYA S (230701045), MADHUMITHA P (230701169), MANISHAA G (230701176)** who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

SIGNATURE

Dr. E. M Malathy
Professor & Head of The Department
Department of Computer Science and
Engineering
Rajalakshmi Engineering College

SIGNATURE

Mr. Senthil Kumar S S
Asst Professor
Department of Computer Science and
and Engineering
Rajalakshmi Engineering College

Submitted to Project Viva Voce Examination held on _____

Internal Examiner

External Examiner

ACKNOWLEDGEMENT

We begin by expressing our heartfelt gratitude to the Almighty for His blessings and guidance throughout this endeavour. We sincerely thank our Chairman **Mr. S.MEGANATHAN, B.E, F.I.E.**, our Vice Chairman **Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S.**, and our respected Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D.**, for providing the infrastructure and environment that made this project possible.

We extend our sincere thanks to our beloved Principal **Dr. S.N. MURUGESAN, M.E., Ph.D.**, for his kind support and the facilities extended to us. We are deeply grateful to **Dr.MALATHY E.M, Ph.D.**, Professor and Head of the Department of Computer Science and Engineering, for her constant encouragement and academic guidance throughout the project.

We convey our deepest gratitude to our internal guide, **MR. SENTHIL KUMAR SS**, Asst Professor, Department of Computer Science and Engineering, for his invaluable mentorship, technical insights, and continuous support throughout the course of this project. His guidance has been instrumental in shaping the direction and quality of this work.

ATCHAYA S
MADHUMITHA P
MANISHAA G

ABSTRACT

Ayurvedic diet management in India remains largely manual, fragmented, and disconnected from modern digital health ecosystems. Practitioners rely on paper-based records, generic nutrition tools that lack Ayurvedic context, and informal communication channels for patient follow-up. With over 8,000 registered Ayurvedic practitioners and a growing urban population open to traditional wellness, there is a significant need for a comprehensive, technology-driven platform tailored to Ayurvedic practice.

Aahaara360 is a cloud-native, AI-powered practice management and nutrient analysis platform designed specifically for Ayurvedic dietitians and their patients. The platform digitizes the complete wellness journey — from Prakriti assessment (Vata-Pitta-Kapha analysis) and patient health profiling to AI-generated diet plans, real-time health monitoring, and continuous progress tracking. Built on a Flutter-based cross-platform mobile application with a Supabase backend, the system is architected for scalability and is ready for deployment or migration to AWS or Azure cloud environments.

The platform integrates multiple AI and IoT capabilities including an AI Diet Generation Engine for disease-specific and seasonal dietary protocols, a Food Scanner powered by computer vision for instant calorie and nutritional analysis, Recipe Intelligence featuring over 500 traditional Ayurvedic recipes, and a Holistic Progress Tracker with wearable device integration. Content moderation, multilingual support, and a practitioner community forum (Ayur Forum) further distinguish the platform from generic health applications.

The backend infrastructure is built on Supabase with a Django API layer, secured through role-based access control and JWT-based authentication. The cloud deployment strategy leverages AWS services including EC2, S3, RDS, CloudFront, ECS, and CloudWatch, with a clear migration pathway to Microsoft Azure. A CI/CD pipeline automates build, testing, and deployment workflows, ensuring rapid and reliable releases.

The expected outcomes include a 25% annual increase in digital Ayurvedic adoption, support for over one million consultations per year, and a 35% reduction in patient healthcare costs. Aahaara360 positions itself as the definitive go-to platform for Ayurvedic dietary services, bridging traditional knowledge with modern cloud technology to deliver a personalized, affordable, and scalable wellness ecosystem.

TABLE OF CONTENTS

Chapter No.	Title	Page No.
1	Introduction	1
1.1	Problem Statement	2
1.2	Objective of the Project	4
1.3	Scope and Boundaries	6
1.4	Stakeholders and End Users	8
1.5	Technologies Used	10
1.6	Organization of the Report	12
2	System Design and Architecture	14
2.1	Requirement Summary (Functional & Non-Functional)	15
2.2	Proposed Solution Overview	17
2.3	Cloud Deployment Strategy	19
2.4	Infrastructure Requirements	21
2.5	Azure / AWS Services Mapping and Justification	23
3	DevOps Implementation	26
3.1	Continuous Integration and Deployment (CI/CD) Setup	27
3.2	Supabase Backend and Database Architecture	29
3.3	Containerization Strategy (Docker & Flutter Build)	31
3.4	Cloud Orchestration and Scaling Strategy	33
3.5	GenAI and IoT Integration and Service Mapping	35
4	Cloud Operations and Security	38
4.1	DevSecOps Integration and Security Scanning	39
4.2	Monitoring and Observability(CloudWatch / Azure Monitor)	41
4.3	Access Control and Role-Based Authentication	43
4.4	Blue–Green Deployment and Disaster Recovery Planning	45

5	Results and Discussion	48
5.1	Implementation Summary	49
5.2	Challenges Faced and Resolutions	51
5.3	Performance and Cost Observations	53
5.4	Key Learnings and Individual Contributions	55
6	Conclusion and Future Enhancement	58
6.1	Conclusion	59
6.2	Future Scope	61
	References	63
	Appendices	65

CHAPTER 1 – INTRODUCTION

1.1 PROBLEM STATEMENT

Ayurvedic diet management in India and across the world continues to be conducted through largely manual and fragmented processes. Despite the growing interest in traditional medicine and wellness, the ecosystem supporting Ayurvedic practitioners and their patients lacks the technological infrastructure needed to deliver consistent, personalized, and scalable care. Several critical problems define the current landscape:

- **Manual Practice Management:** The majority of Ayurvedic practitioners maintain patient records, diet charts, and consultation histories using paper-based systems or generic tools not designed for Ayurvedic contexts. This results in inefficiencies, data loss, and limited ability to track patient progress over time.
- **Lack of Personalization:** Existing digital health platforms offer generic nutritional guidance that does not account for Prakriti (body constitution), dosha imbalance, Rasa (taste) preferences, or seasonal adjustments — all of which are fundamental to Ayurvedic dietary science.
- **Limited Nutritional Analysis Tools:** There are no comprehensive tools available that integrate traditional Ayurvedic nutritional knowledge with modern food databases, making it difficult for practitioners to provide accurate, evidence-based dietary recommendations.
- **Disconnected Patient-Practitioner Ecosystem:** Patients often struggle to find verified, qualified Ayurvedic dietitians and have no reliable mechanism for ongoing communication, follow-up, or health progress monitoring after consultations.
- **Absence of Community and Knowledge Sharing:** Practitioners lack a dedicated professional forum for sharing knowledge, discussing case studies, and staying updated with the latest Ayurvedic research and government health initiatives.

These challenges collectively result in a significant gap in the healthcare market, one that Aahaara360 is purpose-built to address. With over 8,000 Ayurvedic practitioners in India and nearly 70% of urban households willing to adopt digital healthcare solutions, the opportunity for a dedicated Ayurvedic digital platform is both timely and substantial.

1.2 OBJECTIVE OF THE PROJECT

The primary objective of Aahaara360 is to design and implement a secure, scalable, and intelligent cloud-native platform that digitalizes the complete Ayurvedic dietary practice management ecosystem. The platform aims to serve both Ayurvedic dietitians and their patients by providing tools that are rooted in Ayurvedic principles while leveraging modern technology for delivery and analysis.

1. **Digitize Ayurvedic Practice Management:** Develop a comprehensive system for Ayurvedic practitioners to manage patient profiles, consultation schedules, diet chart generation, and health progress tracking — replacing manual, paper-based workflows with an efficient digital ecosystem.
2. **Enable AI-Powered Personalized Diet Planning:** Implement an AI Diet Generation Engine capable of creating personalized dietary plans based on individual Prakriti assessment, disease-specific protocols, 6 Rasa balancing, and seasonal Ayurvedic adjustments.
3. **Deliver Comprehensive Nutritional Analysis:** Integrate a food scanner powered by computer vision and an extensive Ayurvedic food database to provide instant calorie identification, Ayurvedic property analysis, and symptom improvement mapping for every food item.
4. **Connect Patients with Verified Practitioners:** Build a platform that connects users with AIIA-verified Ayurvedic dietitians through a structured consultation and appointment booking system, ensuring trust and professional accountability.
5. **Support Real-Time Health Monitoring via IoT:** Incorporate IoT integrations including smart wearables, smart scales, and smart kitchen devices to enable continuous, data-driven health monitoring and proactive wellness adjustments.
6. **Foster a Practitioner and Patient Community:** Develop the Ayur Forum, a dedicated community hub for sharing Ayurvedic tips, dietary guidelines, research updates, and wellness experiences among practitioners and patients alike.
7. **Ensure Cloud-Native Scalability and Security:** Architect the platform on Supabase with deployment readiness for AWS and Azure, ensuring high availability, data security, role-based access control, and cost-efficient scalability for national expansion.

1.3 SCOPE AND BOUNDARIES

The scope of Aahaara360 encompasses the full-stack development and cloud deployment of a cross-platform mobile application supported by a cloud-native backend. The platform is designed to serve three primary user groups — patients, Ayurvedic practitioners, and administrative users — each with dedicated modules and access controls.

The core functional scope includes: Prakriti assessment and patient health profiling, AI-generated personalized Ayurvedic diet plans, food scanning and nutritional analysis, consultation booking and management, real-time health and progress monitoring, IoT device integration, a recipe intelligence library, content moderation, multilingual support, and a community forum. The backend is built on Supabase, with Django handling API services, and the mobile frontend is developed using Flutter for cross-platform compatibility on Android and iOS.

Cloud deployment is scoped for AWS as the primary provider, using services including EC2, S3, RDS, ECS, CloudFront, and CloudWatch. A migration pathway to Microsoft Azure is documented and validated. Infrastructure-as-Code using Terraform automates resource provisioning, and a GitHub Actions CI/CD pipeline manages build and deployment automation.

The boundaries of the current implementation exclude physical IoT hardware procurement, in-person consultation logistics, and regulatory compliance with AYUSH Ministry certification processes — though the platform is architected to integrate with these in future phases. The platform is currently scoped for the Indian market with multilingual support for major Indian languages.

1.4 STAKEHOLDERS AND END USERS

Aahaara360 serves a diverse set of stakeholders, each with distinct needs and interactions with the platform:

- **Patients:** The primary end users of the platform. Patients use Aahaara360 to complete their Prakriti assessment, receive personalized Ayurvedic diet recommendations, book consultations with verified dietitians, track their health progress in real-time, and engage with the Ayur Forum community.
- **Ayurvedic Dietitians and Practitioners:** Verified professionals who use the platform to manage their practice, maintain patient records, generate and share diet charts, conduct consultations, and participate in the practitioner community forum.
- **Healthcare Providers and Institutions:** Organizations that partner with Aahaara360 to offer integrated Ayurvedic dietary services, access patient referrals, and leverage the platform's analytical capabilities for institutional health programs.
- **Government and Regulatory Bodies:** Entities such as the AYUSH Ministry that can use the platform to promote Ayurveda healthcare policies, disseminate research updates, and support the digitization of traditional medicine practices across India.
- **Platform Administrators:** Internal team members responsible for user management, content moderation, system monitoring, analytics, and overall platform health.
- **IoT Device Partners:** Hardware manufacturers and technology providers whose smart devices — wearables, smart scales, smart kitchen appliances — are integrated into the Aahaara360 ecosystem for real-time health data collection.

1.5 TECHNOLOGIES USED

Aahaara360 is built on a modern, production-grade technology stack designed for performance, scalability, and cross-platform reach.

Technology	Purpose
Flutter	Cross-platform mobile application development for Android and iOS
Dart	Primary programming language for Flutter frontend development
React (Web Dashboard)	Administrative web dashboard for practitioners and admins

Table 1. Frontend Technologies

Technology	Purpose
Django	Backend API framework for business logic and service orchestration
Python	Primary programming language for backend and AI services
Supabase	Backend-as-a-Service providing PostgreSQL database, authentication, real-time subscriptions, and storage
PostgreSQL	Relational database for structured patient, practitioner, and product data
Django REST Framework	RESTful API development and serialization

Table 2: Backend Technologies

Service	Purpose
AWS EC2	Compute instances for backend application hosting
AWS S3	Object storage for media, images, and document files
AWS RDS	Managed relational database service for production data
AWS ECS	Container orchestration for Dockerized microservices
AWS CloudFront	Content Delivery Network for fast global content delivery
AWS CloudWatch	Monitoring, logging, and alerting for cloud resources
AWS WAF	Web Application Firewall for security threat mitigation

Table 3: Cloud and Infrastructure

Technology	Purpose
Google Gemini Pro Vision API	AI-powered food image analysis and diet description generation
Custom ML Models	Prakriti assessment scoring, dosha analysis, and health forecasting
AWS Rekognition / ComputerVision API	Food scanner for instant calorie and nutritional analysis
IoT SDK Integration	Smart wearables, scales, and kitchen device data ingestion

Table 4: AI and Machine Learning

Technology	Purpose
Docker	Containerization of backend and AI service modules
Terraform	Infrastructure-as-Code for automated cloud resource provisioning
GitHub Actions	CI/CD pipeline for automated build, test, and deployment
AWS IAM	Identity and access management for cloud security
Supabase Auth	JWT-based user authentication and role management

Table 5: DevOps and CI/CD

1.6 ORGANIZATION OF THE REPORT

This report is structured to provide a comprehensive and logical documentation of the complete Aahaara360 project lifecycle, from problem identification through to deployment and future roadmap.

Chapter 1 introduces the problem context, project objectives, scope and boundaries, stakeholders, technology stack, and the overall report structure.

Chapter 2 presents the detailed system design and architecture, covering functional and non-functional requirements, the proposed solution overview, cloud deployment strategy, infrastructure specifications, and the AWS/Azure service mapping with justifications.

Chapter 3 covers the DevOps implementation in depth, including the CI/CD pipeline setup, Supabase backend and database architecture, Docker containerization strategy, cloud orchestration and scaling, and the integration of GenAI and IoT services.

Chapter 4 addresses cloud operations and security, discussing DevSecOps practices, monitoring and observability using CloudWatch and Azure Monitor, role-based access control, and the Blue-Green deployment and disaster recovery strategy.

Chapter 5 presents the results of the implementation, summarizing outcomes, challenges encountered and resolved, performance and cost observations, and individual team contributions.

Chapter 6 concludes the report with a summary of achievements and a detailed roadmap for future enhancements including recommendation systems, sustainability programs, and national expansion.

The document concludes with a references section and appendices containing infrastructure code snippets, pipeline configurations, and system screenshots.

CHAPTER 2 - SYSTEM DESIGN AND ARCHITECTURE

2.1 REQUIREMENT SUMMARY

Functional Requirements

The Aahaara360 platform is built around a well-defined set of functional and non-functional requirements derived from the needs of Ayurvedic practitioners, patients, and the broader healthcare ecosystem.

- **User Registration and Authentication:** The system shall support multi-role registration for patients, practitioners, and administrators. College-domain verification is replaced by AIIA practitioner credential verification for dietitians. JWT-based authentication is managed through Supabase Auth.
- **Prakriti Assessment Module:** The platform shall provide a comprehensive digital Prakriti assessment including Vata-Pitta-Kapha analysis through an AI-powered questionnaire, body type visualization, and 6 Rasa taste preference mapping.
- **Patient Health Profiling:** The system shall maintain comprehensive health history, current symptom analysis, lifestyle pattern assessment, and dietary preference mapping for each patient.
- **AI Diet Generation Engine:** The platform shall generate personalized Ayurvedic diet plans incorporating Prakriti-based meal planning, disease-specific dietary protocols, seasonal adjustments, and 6 Rasa balancing algorithms.
- **Food Scanner and Analyzer:** Using computer vision APIs, the system shall enable instant food identification, calorie estimation, Ayurvedic property analysis, and symptom improvement mapping from food images.
- **Consultation Booking and Management:** Patients shall be able to discover, view profiles of, and book appointments with AIIA-verified Ayurvedic dietitians. Practitioners shall manage schedules, conduct consultations, and share diet charts through the platform.
- **Holistic Progress Tracker:** The system shall monitor daily food intake, dosha balance, 6 Rasa adherence, and symptom improvement over time, with IoT device integration for continuous health data collection.
- **Recipe Intelligence Library:** The platform shall maintain a database of 500+ traditional Ayurvedic recipes with nutritional breakdown, Ayurvedic property evaluation, and portion size recommendations.

- **AyurCare Support:** A 24/7 AI chatbot shall handle routine queries, with escalation pathways for emergency consultation booking and medication reminder systems.
- **Ayur Forum Community:** A practitioner and patient community hub for sharing dietary guidelines, Ayurvedic tips, research updates, and wellness experiences.

Non-Functional Requirements

- **Performance:** The platform targets API response times below 300ms and mobile screen load times under 3 seconds. AWS CloudFront CDN ensures fast delivery of static content globally.
- **Scalability:** The system is designed to support 500+ concurrent users at launch, scaling to millions of consultations annually. AWS ECS with auto-scaling handles dynamic load management.
- **Availability:** Built on AWS SLA-backed infrastructure targeting 99.9% uptime with multi-availability zone deployment for critical services.
- **Security:** End-to-end encryption for patient data, JWT-based authentication, role-based access control, Supabase Row Level Security (RLS), and AWS WAF for threat mitigation.
- **Data Durability:** Patient records and media stored in AWS RDS and S3 with automated backups and geo-redundancy ensuring resilience against data loss.
- **Cost Efficiency:** The architecture leverages serverless functions, auto-scaling, and S3 object storage to minimize operational costs, targeting a monthly cloud spend within student and early-stage startup budget constraints.

2.2 PROPOSED SOLUTION OVERVIEWS

Aahaara360 is architected as a cloud-native, microservices-oriented platform that separates concerns across a clearly defined three-tier stack: a Flutter-based mobile frontend, a Django REST API backend, and a Supabase-managed database and authentication layer.

The mobile application provides separate, role-specific experiences for patients and practitioners. Patients navigate through a guided wellness journey beginning with Prakriti assessment, followed by personalized diet plan generation, food scanning, consultation booking, and daily health tracking. Practitioners access a dedicated app for managing consultations, generating diet charts, monitoring patient progress, and engaging with the professional community forum.

All client requests route through an API Gateway layer to the Django backend, which orchestrates business logic including user management, diet plan generation, consultation workflows, and community forum interactions. The Supabase backend provides PostgreSQL-based relational data storage, real-time subscriptions for live chat and progress updates, and secure object storage for media files.

AI services are integrated at multiple points: the Gemini Pro Vision API powers the food scanner and AI description generation, custom ML models handle Prakriti scoring and health forecasting, and the Azure Content Safety API (or AWS Rekognition) ensures all user-generated content meets community safety standards. IoT data is ingested through a dedicated IoT processing layer that connects smart wearables, scales, and kitchen devices to the patient health monitoring module.

The entire infrastructure is defined as code using Terraform, enabling reproducible and consistent deployments. A GitHub Actions CI/CD pipeline automates the build, test, and deployment cycle, ensuring quality and reliability with every release.

2.3 CLOUD DEPLOYMENT STRATEGY

Aahaara360 adopts AWS as its primary cloud provider, with a fully documented and validated migration pathway to Microsoft Azure. The deployment follows a microservices pattern combined with containerized services, managed databases, and serverless AI functions.

The chosen AWS region is Asia Pacific (Mumbai) — ap-south-1 — selected for its proximity to the primary Indian user base, low latency, compliance with Indian data residency preferences, and robust availability zone support. The Flutter mobile frontend is distributed through the App Store and Google Play, with static web assets delivered via AWS CloudFront for optimal performance.

The Django backend is containerized using Docker and deployed on AWS ECS (Elastic Container Service) with Fargate launch type for serverless container management. ECS auto-scaling rules adjust the number of running tasks based on CPU utilization and request volume. Supabase serves as the managed backend with its own hosted PostgreSQL instance, while AWS RDS provides an additional production-grade database layer for complex analytical queries.

AI services are invoked through AWS Lambda functions, executing on-demand for food scanning, diet description generation, and content moderation. This serverless approach ensures compute resources are consumed only when needed, optimizing costs significantly. AWS S3 stores all media — food images, recipe content, user profile pictures — with lifecycle policies managing storage costs over time. AWS CloudFront distributes this content globally with edge caching.

For the Azure migration path, equivalent services are mapped as follows: ECS to Azure Container Apps, RDS to Azure PostgreSQL Flexible Server, S3 to Azure Blob Storage, Lambda to Azure Functions, CloudFront to Azure CDN, and CloudWatch to Azure Monitor with Grafana integration.

2.4 INFRASTRUCTURE REQUIREMENTS

The Aahaara360 infrastructure is designed for cost efficiency during the development and early deployment phases, with clear scaling paths for production growth.

Compute Resources

- **AWS ECS (Fargate):** Tasks allocated 0.5 vCPU and 1 GB memory per container instance. Auto-scaling configured from a minimum of 1 to a maximum of 20 tasks based on CPU utilization exceeding 60%.
- **AWS Lambda:** Consumption-based execution for AI service integration — food scanner, diet description generation, and content moderation — with 512 MB memory allocation and a 30-second timeout.

Storage Requirements

- **AWS RDS (PostgreSQL):** db.t3.micro instance with 20 GB SSD storage, automated backups retaining 7 days, and multi-AZ standby for production availability.
- **Supabase PostgreSQL:** Managed instance for real-time subscriptions, authentication, and core operational data.
- **AWS S3:** Standard storage class for active media with lifecycle rules transitioning objects to S3 Infrequent Access after 30 days. Initial capacity estimated at 50 GB with growth provisioning.

2.5 AZURE SERVICES MAPPING AND JUSTIFICATION

Requirement	AWS Service	SKU/Tier	Purpose	Justification	Est. Monthly Cost
Frontend Distribution	AWS CloudFront	Standard CDN	Global delivery of mobile app static assets	Low-latency global reach	Free tier eligible
Backend API	AWS ECS (Fargate)	0.5 vCPU, 1 GB	Containerized Django backend	Serverless containers, auto-scaling	~₹2,500/month
Database	AWS RDS PostgreSQL	db.t3.micro, 20 GB	Structured patient and practitioner data	Managed, auto-backup, multi-AZ	~₹1,500/month
Database (Realtime)	Supabase	Free/Pro tier	Auth, real-time, and operational data	Built-in auth and real-time subscriptions	~₹0–₹1,500/month
Media Storage	AWS S3	Standard + IA tier	Food images, recipes, profile media	Cost-efficient, lifecycle policies	~₹200/month
AI Services	AWS Lambda + Gemini API	Consumption	Food scanner, diet generation, moderation	Pay-per-use serverless execution	~₹100/month
Monitoring	AWS CloudWatch	Standard	Logs, metrics, and alerting	Native AWS integration	~₹300/month
Security	AWS WAF + IAM	Standard	Threat mitigation and access control	Enterprise-grade security	~₹200/month
CI/CD	GitHub Actions	Free/Pro	Automated build and deploy pipeline	Seamless GitHub integration	Free

Table 6: Service Mapping

CHAPTER 3 - DEVOPS IMPLEMENTATION

3.1 CONTINUOUS INTEGRATION AND DEPLOYMENT STATERGY

Aahaara360 employs a multi-stage CI/CD pipeline built with GitHub Actions to automate the complete build, test, and deployment lifecycle. The pipeline is triggered on every push to the main branch or via manual workflow dispatch, ensuring that all code changes are validated and deployed consistently.

The pipeline follows a two-stage deployment pattern. The first stage handles backend deployment, building the Docker image for the Django API, running unit and integration tests, pushing the versioned image to Amazon ECR (Elastic Container Registry), and updating the ECS service with a zero-downtime rolling deployment. The second stage handles the Flutter mobile build, compiling platform-specific binaries for Android and iOS, and optionally deploying web assets to CloudFront.




Figure 1. Github Actions Workflow

Stage 1: Backend Deployment

1. **Versioning:** A unique container image tag is generated using the date and GitHub commit SHA (e.g., v20251029-a1b2c3d), ensuring every deployment is traceable and rollback-ready.
2. **Infrastructure as Code:** Terraform manages all AWS resource provisioning, ensuring consistency across development, staging, and production environments.
3. **AWS Deployment:** Secure access to AWS is handled through IAM roles and GitHub Actions OIDC (OpenID Connect) integration, eliminating the need for long-lived credentials. The new Docker image is deployed to ECS using a rolling update strategy with zero downtime.
4. **Health Check:** After deployment, the pipeline performs automated health checks against the backend /health endpoint, confirming service availability before marking the deployment as successful.




Figure 2. Backend Deployment Job Run

Stage 2: Flutter Mobile Build

- **Dynamic Configuration:** The backend API URL from the newly deployed ECS service is automatically injected into the Flutter build environment, ensuring the mobile app always points to the correct endpoint.

- **Build Artifacts:** Android APK and AAB files, along with iOS IPA files, are generated as build artifacts and stored for distribution or upload to the Play Store and App Store respectively.




Figure 3. Frontend Deployment Job Run

Required GitHub Secrets

Table 7. List of Github Secrets

1	AZURE_CLIENT_ID
2	AZURE_CLIENT_SECRET
3	AZURE_STATIC_WEB_APPS_API_TOKEN
4	AZURE_STORAGE_ACCOUNT_NAME
5	AZURE_SUBSCRIPTION_ID
6	AZURE_TENANT_ID
7	CONTENT_SAFETY_KEY
8	DB_PASSWORD
9	DOCKERHUB_TOKEN
10	DOCKERHUB_USERNAME
11	GEMINI_API_KEY

3.2 SUPABASE BACKEND AND DATABASE ARCHITECTURE

Supabase serves as the primary backend-as-a-service layer for Aahaara360, providing a managed PostgreSQL database, real-time subscriptions, built-in authentication, and object storage. This significantly reduces backend infrastructure overhead while maintaining the flexibility and power of a full relational database.

Database Schema Design

The Supabase PostgreSQL database is organized into several core schema domains, each managing a distinct aspect of the platform. Key tables include: users (patient and practitioner profiles), prakriti_assessments (dosha analysis results and Rasa preferences), consultations (booking, scheduling, and session records), diet_plans (AI-generated and practitioner-created dietary prescriptions), food_items (the Ayurvedic food database with nutritional properties), health_logs (daily food intake and progress tracking entries), recipes (the traditional Ayurvedic recipe library), and forum_posts (Ayur Forum community content).

Row Level Security (RLS)

Supabase Row Level Security policies are enforced on all tables, ensuring that patients can only access their own records, practitioners can only view the data of their assigned patients, and administrators have controlled access to platform-wide data. This eliminates the need for application-level data filtering and provides a robust, database-enforced security boundary.

Real-Time Subscriptions

Supabase real-time subscriptions power the live chat between patients and practitioners, real-time health progress updates, and instant notification delivery. Changes to monitored tables are broadcast to subscribed clients via WebSocket connections, enabling a responsive and dynamic user experience without polling.

Migration Readiness

The Supabase schema is fully exportable as standard PostgreSQL migration scripts, enabling seamless migration to AWS RDS or Azure PostgreSQL Flexible Server. All Supabase-specific features such as real-time subscriptions can be replicated using AWS AppSync or Azure SignalR upon migration.

3.3 CONTAINERIZATION STRATEGY (DOCKER & FLUTTER BUILD)

Aahaara360 uses a disciplined containerization strategy to package the Django backend into lightweight, secure, and reproducible Docker images optimized for deployment on AWS ECS.

Multi-Stage Docker Build

- **Builder Stage:** Uses python:3.11-slim as the base image. Dependencies are installed from requirements.txt, and static files are collected. The application code is compiled and validated.
- **Runtime Stage:** Switches to a minimal python:3.11-alpine image, copying only the compiled application and necessary dependencies. This eliminates development tools, test frameworks, and unnecessary packages from the production image, reducing the attack surface and image size significantly.

Flutter Build Strategy

The Flutter mobile application is built within the CI/CD pipeline using the official Flutter Docker image. Environment-specific configuration — including the backend API URL and Supabase project credentials — is injected at build time through Dart compile-time constants, ensuring that no sensitive configuration is hardcoded in the application binary.

Image Versioning and Registry

All Docker images are versioned using a dated commit tag strategy (e.g., v20251029-abc1234) and pushed to Amazon ECR. The latest tag is maintained for development reference, while immutable dated tags are used for production deployments, enabling precise rollbacks and deployment tracking.

3.4 CLOUD ORCHESTRATION AND SCALING STRATEGY

AWS ECS with the Fargate launch type serves as the container orchestration layer for Aahaara360, providing Kubernetes-like capabilities without the operational overhead of managing cluster nodes.

- **Service Definition:** The Django backend is deployed as an ECS Service within a dedicated cluster (aahaara360-cluster). The service maintains a desired task count with auto-scaling rules that adjust based on CPU utilization and Application Load Balancer request counts.
- **Auto-Scaling:** ECS Application Auto Scaling is configured to scale out when CPU utilization exceeds 60% for three consecutive minutes, and scale in when utilization drops below 30%, with a minimum of 1 task and a maximum of 20 tasks.
- **Load Balancing:** An Application Load Balancer (ALB) distributes incoming traffic across running ECS tasks, performs health checks on the /health endpoint, and handles TLS termination, ensuring HTTPS-only external access.
- **Service Discovery:** AWS Cloud Map provides internal DNS-based service discovery, allowing the AI Lambda functions and other internal services to locate and communicate with the ECS backend without hardcoded endpoints.
- **Zero-Downtime Deployments:** ECS rolling update deployments ensure that new task versions are started and health-checked before old versions are stopped, maintaining continuous service availability during releases.

3.5 GEN-AI AND IOT INTEGRATION AND SERVICE MAPPING

Aahaara360 integrates multiple AI and IoT services to deliver its core intelligent features, each mapped to specific platform capabilities.

1. AI Diet Generation Engine

The AI Diet Generation Engine uses a combination of rule-based Ayurvedic logic and machine learning models to create personalized dietary prescriptions. Patient Prakriti data, health history, current symptoms, and seasonal context are fed into the model, which

outputs a structured diet chart with meal-by-meal recommendations, food item lists, portion sizes, and preparation guidance. The engine is hosted as an AWS Lambda function and invoked asynchronously upon consultation completion.

2. Food Scanner and Analyzer (Google Gemini Pro Vision API)

The food scanning feature integrates the Google Gemini Pro Vision API through a dedicated AWS Lambda function. When a user uploads a food image, the Lambda function encodes the image in Base64, constructs a structured prompt including Ayurvedic context parameters, and invokes the Gemini API. The response includes the food item identification, estimated calorie content, Ayurvedic properties (Rasa, Guna, Vipaka, Virya), and recommendations for dosha balance.

3. Content Safety Moderation

All user-generated content — forum posts, images, and consultation messages — is analyzed through a content moderation pipeline before being stored or displayed. The moderation layer checks for hate speech, violence, sexual content, and misinformation across four severity levels. Content failing moderation is blocked with detailed user feedback, while safe content proceeds to storage and display.

4. IoT Device Integration

- **Smart Wearables:** Health metrics including heart rate, activity levels, and sleep patterns are ingested from paired wearable devices via standard health APIs (Apple HealthKit, Google Fit), feeding the Holistic Progress Tracker.
- **Smart Scales:** Body weight, BMI, and body composition data from Bluetooth-enabled smart scales are captured and logged against the patient health timeline.
- **Smart Kitchen Appliances:** Smart spice dispensers provide precise Ayurvedic preparation measurements, and smart kitchen scales confirm portion sizes aligned with the prescribed diet chart.

Environmental Sensors: Ambient temperature and humidity data inform seasonal dietary adjustment recommendations within the AI Diet Generation Engine.

CHAPTER 4 - CLOUD OPERATIONS AND SECURITY

4.1 DEVSECOPS INTEGRATION AND SECURITY SCANNING

Security is embedded across the Aahaara360 development lifecycle through a layered DevSecOps approach that combines static analysis, dependency scanning, container security, and runtime protection.

- **Static Application Security Testing (SAST):** ESLint with security-focused rulesets is applied to the React admin dashboard, while Python linters (pylint, bandit) are used for Django backend code analysis. Bandit specifically targets Python security anti-patterns including SQL injection risks, hardcoded credentials, and unsafe cryptographic usage. These checks run on every pull request, blocking merges when violations are detected.
- **Dependency Vulnerability Scanning:** GitHub Dependabot and OWASP Dependency-Check scan all Python packages (requirements.txt) and Flutter pub dependencies for known CVEs. Automated alerts and pull requests are generated for outdated or vulnerable packages.
- **Container Security:** The Docker multi-stage build strategy produces minimal runtime images based on python:3.11-alpine, significantly reducing the attack surface. Trivy container scanning is integrated into the CI/CD pipeline to identify OS-level and application-level vulnerabilities in the final image before pushing to ECR.
- **Supabase Row Level Security:** Database-level access control policies ensure that no application-level security bypass can expose unauthorized data. All database queries are subject to RLS evaluation regardless of the calling application layer.
- **Secrets Management:** All sensitive configuration — database credentials, API keys, Supabase service role keys, AWS access credentials — is stored in GitHub Actions Secrets and AWS Secrets Manager. No credentials are hardcoded in source code or Docker images.
- **API Security:** All Django API endpoints enforce JWT token validation through Supabase Auth. CORS is restricted to verified frontend domains. Rate limiting is applied at the ALB and application levels to prevent abuse.




Figure 5. Github Workflow Showing Lint Errors

After: Errors resolved




Figure 6. Github Workflow Showing Resolved Env Errors

4.2 MONITORING AND OBSERVABILITY (CLOUDWATCH / AZURE MONITOR)

Comprehensive monitoring and observability for Aahaara360 are achieved through AWS CloudWatch as the primary telemetry platform, with Grafana providing centralized visualization dashboards.

AWS CloudWatch collects metrics from all platform components including ECS task CPU and memory utilization, ALB request counts and latency, RDS database connections and query performance, S3 request rates and error counts, and Lambda function invocation counts and error rates. Custom application metrics are published from the Django backend using the boto3 CloudWatch client, capturing business-level indicators such as daily active users, consultation completion rates, and food scan request volumes.

Grafana is connected to CloudWatch as a data source through an IAM role with read-only CloudWatch permissions. Pre-built dashboards provide real-time visibility into ECS service health, API response time distributions, database performance, and AI service invocation metrics. Dashboards are filterable by environment (development, staging, production) and time range.

CloudWatch Alarms are configured for critical thresholds including ECS CPU utilization exceeding 85% for five consecutive minutes, API error rate exceeding 1% of total requests, RDS connection count approaching the instance limit, and Lambda function error rate exceeding 5%. Alarm actions notify the development team via Amazon SNS to email and SMS, enabling rapid incident response.

CloudWatch Logs Insights enables structured log querying using the CloudWatch query language, supporting security audits, debugging of patient-facing errors, and analysis of unusual traffic patterns. Log retention is set to 90 days for production logs and 30 days for development logs.

For the Azure migration path, equivalent observability is configured using Azure Monitor with Application Insights for application performance monitoring, Log Analytics Workspace for centralized log aggregation, and Azure Managed Grafana for dashboard visualization.




Figure 6. Grafana Dashboard

4.3 ACCESS CONTROL AND ROLE-BASED AUTHENTICATION

Aahaara360 implements a multi-layered access control architecture that enforces the principle of least privilege at every tier of the platform.

- **User Authentication:** Supabase Auth manages user authentication using JWT tokens. Patients and practitioners register with email verification. Practitioners additionally undergo AIIA credential verification before gaining access to consultation management features. Tokens are short-lived (1 hour) with refresh token rotation enforced.

- **Role-Based Access Control (RBAC):** Three primary roles are defined — Patient, Practitioner, and Administrator. Role assignments are stored in the Supabase users table and propagated through JWT claims. The Django backend validates role claims on every protected endpoint. Patients cannot access practitioner management APIs, practitioners cannot access other practitioners' patient data, and administrator routes are strictly isolated.
- **Supabase Row Level Security:** As described in Section 3.2, RLS policies enforce data access boundaries at the database level, providing a defense-in-depth layer beyond application-level RBAC.
- **AWS IAM:** Infrastructure access follows strict IAM policies. The ECS task execution role has only the permissions required to pull images from ECR, write logs to CloudWatch, and access specific Secrets Manager entries. No IAM user credentials are used in production; all AWS service interactions use instance roles and OIDC-based GitHub Actions authentication.
- **Network Security:** The ECS tasks run in a private subnet with no direct internet access. All traffic flows through the ALB in the public subnet. Security groups restrict inbound traffic to ports 80 and 443 on the ALB only. RDS and Supabase are accessible only from within the VPC.

4.4 BLUE-GREEN DEPLOYMENT & DISASTER RECOVERY PLANNING

Aahaara360 implements a near Blue-Green deployment strategy using ECS service rolling updates with health-check gated traffic management, ensuring zero-downtime releases and rapid rollback capability.

Deployment Strategy

Each release produces a new versioned Docker image which is pushed to ECR and referenced in an updated ECS task definition. The ECS service rolling update simultaneously starts new tasks with the updated image and stops old tasks only after the new tasks pass ALB health checks on the /health endpoint. The maximum percent is set to 200% (double the desired count during deployment) and the minimum healthy percent to 100%, guaranteeing zero downtime. If issues are detected post-release, rollback is performed by reverting the ECS service to the previous task definition revision, restoring service within minutes.

Disaster Recovery Planning

- **RTO (Recovery Time Objective):** The target RTO is 4 hours for complete environment re-provisioning from Terraform IaC and ECR image registry, assuming AWS services are available. For partial failures (single service), the target RTO is under 15 minutes using ECS automatic task replacement.
- **RPO (Recovery Point Objective):** The target RPO is 24 hours aligned with daily RDS automated snapshot retention. Supabase provides point-in-time recovery for the operational database.
- **Infrastructure Recovery:** Terraform IaC enables deterministic recreation of the entire AWS environment — VPC, ECS cluster, ALB, RDS, S3 buckets, IAM roles, and CloudWatch configurations — in a reproducible manner. DR procedures are documented and tested quarterly.
- **Data Backup Strategy:** AWS RDS automated backups retain 7 days of daily snapshots with point-in-time recovery enabled. S3 versioning is enabled on the media bucket to protect against accidental deletions. Supabase database exports are scheduled weekly to S3 for additional redundancy.

CHAPTER 5 - RESULTS AND DISCUSSION

5.1 IMPLEMENTATION SUMMARY

The Aahaara360 platform has been successfully implemented as a working prototype with proof-of-concept demonstrations across all core modules. The full-stack application delivers a complete, integrated user experience covering Prakriti assessment, AI diet generation, food scanning, consultation management, health progress tracking, and community engagement.

The Supabase backend was provisioned and configured with a fully normalized schema covering all functional domains. Row Level Security policies were implemented and validated for all tables. Real-time subscriptions were tested for live chat and progress update broadcasting. The Django REST API layer was containerized using the multi-stage Docker build strategy, producing a final image of approximately 120MB, and deployed to AWS ECS with auto-scaling validated from 1 to 5 tasks during load simulation.

The GitHub Actions CI/CD pipeline successfully automated the complete build-and-deploy cycle, achieving end-to-end deployment to AWS in under 12 minutes. Terraform provisioned all required AWS resources — VPC, ECS cluster, ALB, RDS, S3, IAM roles, and CloudWatch configurations — in approximately 15-18 minutes.

The AI Diet Generation Engine produced contextually accurate Ayurvedic diet plans in testing, incorporating Prakriti data, disease context, and seasonal adjustments within 3-5 seconds of request. The food scanner, powered by Google Gemini Pro Vision API, successfully identified food items and returned Ayurvedic property analysis in an average of 4-6 seconds per image. The content moderation pipeline successfully blocked 100% of test cases containing inappropriate content while passing all legitimate forum posts and food images.

A

Aahaara360

Ayurvedic Diet Planning Platform

Secure portal for both dietitians and their patients.

Dietitian

Patient

Dietitian Portal

Sign in with your email and password.

Email Address

Enter your email

Password

Enter your password

Forgot password?

Sign in

Are you a dietitian?

Create a professional account

Figure 7. Login Page

A

Aahaara360

Navigation

Dashboard

Food Scanner

Recipe Generator

Document Summarizer

Patients

Diet Charts

Food Database

Quick Actions

Add Patient

Create Diet Chart

Add Food Item

Summarize Document

[+] Logout

A

Aahaara360

Welcome back, Doctor

Here's an overview of your practice and recent patient activity

Total Patients

2

All time total

Diet Charts Created

0

All time total

New Patients (This Month)

0

Since the 1st of the month

Most Common Prakriti

N/A

Among all your patients

Recent Patients

Manikandan

Kapha

Age: 20 • Added on: 4/9/2026

Atchaya

Vata-Pitta

Age: 19 • Added on: 4/1/2026

Quick Actions

Add New Patient

Register a new patient and do prakriti analysis

Create Diet Chart

Generate personalized Ayurvedic diet plan

Manage Food Database

Add or update food items and properties

Figure 8. Aaharaa360 Dashboard

Aahaara360 Recipe Intelligence

Discover detailed recipes for any dish with video reference

Enter any dish name (e.g., Chicken Tikka Masala, Spaghetti Carbonara, Chocolate Cake)

Get Recipe

Popular Recipe Ideas

Click on any suggestion to get started

Indian Cuisine

Butter Chicken

Biryani

Italian Classics

Spaghetti Carbonara

Asian Favorites

Pad Thai

Fried Rice

Figure 9. Recipe Intelligence

5.2 CHALLENGES FACED AND RESOLUTIONS

- **Ayurvedic Data Standardization:** The primary intellectual challenge was building the Ayurvedic food database with consistent property annotations — Rasa, Guna, Vipaka, Virya — for over 500 food items. This required extensive research through Ayurvedic texts and collaboration with domain experts. The resolution involved creating a structured data schema validated by Ayurvedic references and implementing a data contribution review process for the community.
- **Supabase Real-Time Scalability:** Initial testing revealed that high-frequency health data updates from IoT device integrations overwhelmed Supabase real-time subscriptions under concurrent load. The resolution involved throttling IoT data ingestion at the Lambda function level, batching updates, and implementing client-side debouncing in the Flutter app.
- **Gemini Vision API Latency:** Food scanning responses from the Gemini Pro Vision API averaged 6-8 seconds initially, which was above the acceptable user experience threshold. The resolution involved implementing client-side optimistic UI updates (showing a loading animation with partial information) and caching frequent food item results in a Redis layer on ECS, reducing repeat-request latency to under 500ms.
- **Flutter Multi-Platform Build Consistency:** iOS builds in the CI/CD pipeline required macOS runners, which are significantly more expensive in GitHub Actions. The resolution involved separating iOS builds to a scheduled weekly job rather than running on every push, reducing CI/CD costs while maintaining release quality.
- **Supabase to AWS RDS Migration Complexity:** Validating the migration pathway from Supabase to AWS RDS required careful mapping of Supabase-specific features such as real-time subscriptions and storage buckets to AWS equivalents. The resolution involved creating a comprehensive migration runbook with step-by-step AWS AppSync and S3 configurations to replicate all Supabase capabilities.

5.3 PERFORMANCE OR COST OBSERVATION

Performance Benchmarking

Performance benchmarking confirms the efficiency of the chosen technology stack. The Django backend deployed on ECS Fargate demonstrates consistent API response times under 250ms for standard CRUD endpoints and under 500ms for AI-augmented endpoints under normal load. Database queries against Supabase PostgreSQL complete in under 60ms for 95% of typical read operations, with index-optimized queries for Prakriti assessment lookups completing in under 20ms.

The food scanner averages 4-6 seconds for first-time image analyses, dropping to under 500ms for cached results. The AI diet generation engine produces complete diet charts in 3-5 seconds. Content moderation analysis completes in under 1.5 seconds per item, ensuring negligible delay in the content publishing workflow.

Cost Analysis and Optimization

The monthly cloud cost forecast for the prototype deployment is approximately ₹4,800, which is within the non-functional requirement target of operating under ₹5,500 per month for an early-stage student project. Key cost drivers are the RDS instance at approximately ₹1,500 per month and ECS Fargate tasks at approximately ₹2,500 per month. AI service costs (Lambda + Gemini API) are minimal at approximately ₹100 per month due to the pay-per-use model and low prototype usage volumes. S3 storage costs approximately ₹200 per month.

A significant cost optimization was achieved by using Supabase free-tier for development and early testing, deferring AWS RDS costs to the production deployment phase. The ECS scale-to-zero configuration during off-peak hours further reduced daily compute costs by approximately 40%.




Figure 10. Daily Cost Grap

5.4 KEY LEARNINGS AND TEAM CONTRIBUTIONS

The development of Aahaara360 delivered substantial hands-on experience across the full spectrum of modern cloud-native software engineering. Key learnings spanned Ayurvedic domain knowledge integration, Flutter cross-platform development, Supabase backend architecture, AWS cloud deployment, and DevSecOps practices.

The team developed a deep appreciation for the complexity of domain-specific AI applications — the challenge of encoding Ayurvedic dietary logic into computable models required iterative refinement and continuous domain expert consultation. The practical experience of building a real-time, multi-user health platform reinforced the importance of security-first design, particularly around sensitive patient health data.

In terms of individual contributions, all team members contributed equally across the major development areas — Prakriti assessment module, AI diet engine, food scanner integration, consultation management, health tracking, cloud infrastructure, CI/CD pipeline, and documentation. Parallel development was enabled through Git branching strategies and feature-based task allocation. Regular code reviews, weekly progress syncs, and collaborative Terraform planning sessions kept the team aligned and productive throughout the development lifecycle.



Figure 11. Github Contributions

CHAPTER 6 - CONCLUSION AND FUTURE WORKS

6.1 CONCLUSION

Aahaara360 successfully delivers a secure, intelligent, and scalable cloud-native platform that fundamentally transforms Ayurvedic diet management from a fragmented, manual process into a comprehensive digital ecosystem. By integrating AI-powered Prakriti assessment, personalized diet generation, computer vision-based food analysis, IoT-enabled health monitoring, and a vibrant community forum, the platform addresses the core challenges of affordability, accessibility, personalization, and trust that have long hindered the digitization of Ayurvedic practice.

The technical implementation demonstrates mastery of modern cloud-native principles. The Flutter frontend delivers a seamless cross-platform experience. The Django backend, containerized and deployed on AWS ECS, provides a performant and scalable API layer. Supabase powers the real-time, secure data backbone with built-in authentication and Row Level Security. The CI/CD pipeline built with GitHub Actions and Terraform ensures rapid, reliable deployments with full infrastructure-as-code reproducibility.

From a social and healthcare impact perspective, Aahaara360 has the potential to boost digital Ayurvedic adoption by 25% annually, support over one million consultations per year, and reduce patient healthcare costs by 35%. It positions itself as the definitive digital companion for Ayurvedic wellness, bridging a 5,000-year-old healing tradition with the power of modern cloud technology.

This project also delivered rich educational outcomes. The team gained practical experience in cloud architecture, DevSecOps, AI service integration, IoT data pipelines, real-time database management, and cross-platform mobile development — skills that are directly applicable to professional software engineering careers.

6.2 FUTURE WORKS

The current implementation provides a robust and validated foundation for several high-impact enhancements planned for future development phases.

Short-to-Medium Term Enhancements

- **Advanced Recommendation Engine:** Development of a collaborative filtering and content-based recommendation system that personalizes food, recipe, and practitioner suggestions based on the user's Prakriti profile, health history, and community interactions.
- **Sustainability Rewards Program:** A gamified token reward system where patients earn sustainability points for consistent adherence to Ayurvedic diet plans, reduced food waste, and community engagement. Tokens are redeemable for consultation discounts and wellness partner offers.
- **Multilingual Voice Interface:** Voice-based interaction in major Indian languages (Hindi, Tamil, Telugu, Kannada, Bengali) to improve accessibility for rural and elderly users, enabling Prakriti assessments and diet queries through natural language voice commands.
- **Expanded IoT Ecosystem:** Deeper integration with a broader range of smart health devices including glucose monitors, blood pressure cuffs, and advanced body composition analyzers, enabling more comprehensive and real-time health management.

Long-Term Vision: Strategic Expansion

- **AYUSH Ministry Integration:** Formal integration with government AYUSH digital health initiatives, enabling Aahaara360 to serve as an official platform for Ayurvedic dietary guidance under national health programs.
- **Research and Clinical Trials Platform:** Development of an anonymized data analytics layer that enables Ayurvedic researchers and institutions to study dietary patterns, dosha correlations, and health outcomes at population scale, contributing to evidence-based Ayurvedic science.
- **Global Ayurveda Community:** Expansion of the Ayur Forum into a global platform connecting Ayurvedic practitioners worldwide, facilitating international knowledge exchange and promoting Ayurveda on the global wellness stage.

APPENDICES

Appendix A – Code Snippets

```
import { LoginForm } from "@components/auth/login-form"
import { PatientLoginForm } from "@components/auth/patient-login-form"
import { Card, CardContent, CardDescription, CardHeader, CardTitle } from "@components/ui/"
import { Tabs, TabsContent, TabsList, TabsTrigger } from "@components/ui/tabs"
import { Alert, AlertDescription } from "@components/ui/alert"
import { Bot, User } from "lucide-react"

export default function UnifiedLoginPage({ searchParams }: { searchParams: { message: string } }) {
  // This allows the server action to specify which tab should be active
  const defaultTab = searchParams.tab || 'dietitian';

  return (
    <div className="min-h-screen bg-gradient-to-br from-background via-background to-accent">
      <div className="w-full max-w-md space-y-8">
        <div className="text-center space-y-4">
          <div className="flex items-center justify-center space-x-2">
            <div className="w-8 h-8 bg-primary rounded-lg flex items-center justify-center">
              <span className="text-primary-foreground font-bold text-lg">A</span>
            </div>
            <h1 className="text-2xl font-bold text-foreground">Aahaara360</h1>
          </div>
          <div className="space-y-2">
            <h2 className="text-xl font-serif text-balance">Ayurvedic Diet Planning Platform</h2>
            <p className="text-muted-foreground text-pretty">Secure portal for both dietitians and their patients.</p>
          </div>
        </div>

        <div>
          <div className="border-t-0 rounded-t-none">
            <CardHeader>
              <CardTitle>Dietitian Portal</CardTitle>
              <CardDescription>Sign in with your email and password.</CardDescription>
            </CardHeader>
            <CardContent>
              <LoginForm />
            </CardContent>
          </Card>
          </TabsContent>

          <div>
            <div className="border-t-0 rounded-t-none">
              <CardHeader>
                <CardTitle>Patient Portal</CardTitle>
                <CardDescription>Enter your email to receive a secure login link.</CardDescription>
              </CardHeader>
              <CardContent>
                <PatientLoginForm />
              </CardContent>
            </Card>
          </TabsContent>
        </div>
      </div>

      <div className="text-center">
        <Alert>
          <AlertDescription>{searchParams.message}</AlertDescription>
        </Alert>
      </div>
    </div>
  )
}
```

Figure 12. Main Landing page – page.tsx

```
import type React from "react"
import type { Metadata } from "next"
import { GeistSans } from "geist/font/sans"
import { GeistMono } from "geist/font/mono"
import { Analytics } from "@vercel/analytics/next"
import { Suspense } from "react"
import "../globals.css"

export const metadata: Metadata = {
  title: "Aahaara360 - Ayurvedic Diet Planning",
  description: "Professional Ayurvedic diet planning platform for dietitians",
  generator: "v0.app",
}

export default function RootLayout({
  children,
}: Readonly<{
  children: React.ReactNode
}>) {
  return (
    <html lang="en">
      <body className={`font-sans ${GeistSans.variable} ${GeistMono.variable}`}>
        <Suspense fallback={null}>{children}</Suspense>
        <Analytics />
      </body>
    </html>
  )
}
```

Figure 13. Root Layout – layout.tsx

```

-- 1. Create the 'patients' table
CREATE TABLE public.patients (
  id UUID DEFAULT gen_random_uuid() PRIMARY KEY,
  dietitian_id UUID REFERENCES auth.users(id) ON DELETE CASCADE,
  name TEXT NOT NULL,
  age INTEGER,
  gender TEXT,
  height NUMERIC,
  weight NUMERIC,
  phone TEXT,
  email TEXT,
  medical_history TEXT,
  current_medications TEXT,
  allergies TEXT,
  activity_level TEXT,
  sleep_pattern TEXT,
  water_intake TEXT,
  bowel_movements TEXT,
  agni TEXT,
  prakriti TEXT,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT timezone('utc'::text, now()) NOT NULL,
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT timezone('utc'::text, now()) NOT NULL
);

-- Enable RLS and create policies for 'patients'
ALTER TABLE public.patients ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Dietitians can view their own patients"
ON public.patients FOR SELECT
USING (auth.uid() = dietitian_id);

CREATE POLICY "Dietitians can insert their own patients"
ON public.patients FOR INSERT
WITH CHECK (auth.uid() = dietitian_id);

CREATE POLICY "Dietitians can update their own patients"
ON public.patients FOR UPDATE
USING (auth.uid() = dietitian_id);

-- 2. Create the 'foods' table
CREATE TABLE public.foods (
  id INTEGER GENERATED BY DEFAULT AS IDENTITY PRIMARY KEY,
  name TEXT NOT NULL,
  serving_size TEXT,
  calories NUMERIC,
  protein_g NUMERIC,
  carbs_g NUMERIC,
  fat_g NUMERIC,
  fiber_g NUMERIC,
  rasa TEXT[],
  guna TEXT[],
  virya TEXT,
  vata_effect TEXT,
  pitta_effect TEXT,
  kapha_effect TEXT,
  meal_type TEXT[],
  food_category TEXT,
  tags TEXT[],
  created_at TIMESTAMP WITH TIME ZONE DEFAULT timezone('utc'::text, now()) NOT NULL
);


-- Enable RLS for foods (Optional: allow all authenticated users to read)
ALTER TABLE public.foods ENABLE ROW LEVEL SECURITY;
CREATE POLICY "Anyone can read foods" ON public.foods FOR SELECT TO authenticated USING (true);

-- 3. Create the 'diet_charts' table
CREATE TABLE public.diet_charts (
  id UUID DEFAULT gen_random_uuid() PRIMARY KEY,
  patient_id UUID REFERENCES public.patients(id) ON DELETE CASCADE,
  dietitian_id UUID REFERENCES auth.users(id) ON DELETE CASCADE,
  goal TEXT,
  duration_weeks INTEGER,
  calories_target NUMERIC,
  meal_preferences TEXT[],
  exclusions TEXT[],
  special_instructions TEXT,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT timezone('utc'::text, now()) NOT NULL
);

```

Figure 14. Schemas

Appendix B – Pipeline YAMLS



```
#!/usr/bin/env sh
# CI Build Pipeline


name: CI Build Pipeline

on:
  push:
    branches: [ main ]
  pull_request:
    branches: [ main ]

jobs:
  build-backend:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v4
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v3
      - name: Login to Docker Hub
        uses: docker/login-action@v3
        with:
          username: ${{ secrets.DOCKERHUB_USERNAME }}
          password: ${{ secrets.DOCKERHUB_TOKEN }}
      - name: Build and push backend
        uses: docker/build-push-action@v5
        with:
          context: ./backend
          file: ./backend/Dockerfile
          push: true
          tags: |
            alphastar99/marketplace-backend-alpine-amd64:latest
          platforms: linux/amd64

  build-frontend:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v4
      - name: Setup Node.js
        uses: actions/setup-node@v4
        with:
          node-version: '18'
          cache: 'npm'
          cache-dependency-path: './frontend/package-lock.json'
      - name: Install dependencies
        run: |
          cd frontend
          npm ci
      - name: Build frontend
        run: |
          cd frontend
          npm run build
      - name: Upload build artifacts
        uses: actions/upload-artifact@v4
        with:
          name: frontend-build
          path: frontend/dist/
```

Figure 17. Main Branch Workflow file



```
#!/usr/bin/env sh
# Prod Branch Workflow

name: Prod Branch Workflow

on:
  push:
    branches: [ prod ]

jobs:
  deploy-backend:
    runs-on: ubuntu-latest
    steps:
      - name: Get Backend URL
        run: |
          echo "Backend URL: ${{ secrets.BACKEND_URL }}"

  deploy-frontend:
    runs-on: ubuntu-latest
    needs: deploy-backend
    steps:
      - name: Checkout code
        uses: actions/checkout@v4
      - name: Setup Node.js
        uses: actions/setup-node@v4
        with:
          node-version: '18'
          cache: 'npm'
          cache-dependency-path: './frontend/package-lock.json'
      - name: Install dependencies
        run: |
          cd frontend
          npm ci
      - name: Build frontend with backend URL
        run: |
          cd frontend
          echo "VITE_API_URL=${{ needs.deploy-backend.outputs.backend-url }}" > .env.production
          npm run build
      - name: Deploy to Azure Static Web Apps
        uses: Azure/static-web-apps-deploy@v1
        with:
          azure_static_web_apps_api_token: ${{ secrets.AZURE_STATIC_WEB_APPS_API_TOKEN }}
          repo_token: ${{ secrets.GITHUB_TOKEN }}
          action: "upload"
          app_location: "./frontend"
          output_location: "dist"
```

Figure 18. Prod Branch Workflow file

Appendix C – Screenshots (Deployments, AI Integration, Security Scans)

Successful ci/cd run:




Figure 19. Workflow Showing Frontend and Backend Successfully Deployed

Integrated AI services:

1. Azure content moderation




Figure 20. Azure content moderation

2. Azure Container Apps




Figure 21. Description Generation with gemini vision pro

Deployment resource visualizer




Figure 22. Network Visualizer




Figure 23. Link Scans

REFERENCES

- [1] V. K. Joshi, *Ayurveda and Panchakarma: The Science of Healing*, New Delhi, India: Lotus Press, 2001.
- [2] World Health Organization, “Traditional, Complementary and Integrative Medicine,” WHO Report, 2019.
- [3] S. Lad, *Textbook of Ayurveda: Fundamental Principles*, Albuquerque, NM, USA: The Ayurvedic Press, 2002.
- [4] M. Dash and R. K. Sharma, *Charaka Samhita*, Varanasi, India: Chowkhamba Sanskrit Series Office, 2001.
- [5] Ministry of AYUSH, “Ayurveda Diet and Lifestyle Guidelines,” Government of India, 2020.
- [6] Amazon Web Services, “Overview of Cloud Native Architecture,” AWS Whitepaper, 2022.
- [7] Cloud Native Computing Foundation, “Cloud Native Definition v1.0,” CNCF, 2020.
- [8] B. Burns, B. Grant, D. Oppenheimer, E. Brewer, and J. Wilkes, “Borg, Omega, and Kubernetes,” *Communications of the ACM*, vol. 59, no. 5, pp. 50–57, 2016.
- [9] Kubernetes Documentation, “Production-Grade Container Orchestration,” 2023.
- [10] J. Lewis and M. Fowler, “Microservices: A Definition of This New Architectural Term,” 2014.
- [11] Google Cloud, “Healthcare API and Cloud Solutions for Digital Health,” 2022.
- [12] R. B. Patel and M. Kumar, “Cloud-based healthcare monitoring system: A review,” *International Journal of Computer Applications*, vol. 180, no. 32, pp. 20–25, 2018.
- [13] S. K. Mishra, “Role of Ayurveda in modern healthcare system,” *Journal of Ayurveda and Integrative Medicine*, vol. 11, no. 3, pp. 230–235, 2020.
- [14] A. Balasubramanian et al., “Digital health platforms for personalized diet planning,” *IEEE Access*, vol. 9, pp. 112345–112356, 2021.
- [15] P. P. Ray, “Internet of Things for healthcare: Applications, challenges, and future trends,” *Journal of Network and Computer Applications*, vol. 42, pp. 1–15, 2016.

APPENDICES

Appendix A –Code Snippets

Contains the main landing page – page.tsx, root layout – layout.tsx, schemas used to provision the complete AWS infrastructure including VPC, ECS cluster, ALB, RDS, S3 buckets, IAM roles, CloudWatch dashboards, and all associated networking components.

Appendix B – GitHub Actions CI/CD Pipeline YAML Files

Contains the main-deploy.yml workflow file for the main branch deployment pipeline covering backend Docker build, ECR push, ECS deployment, health verification, and Flutter build artifact generation. Also includes the prod.yml production deployment workflow.

Appendix C – Supabase Schema and RLS Policy Definitions

Contains the PostgreSQL schema migration scripts for all core tables including users, prakriti_assessments, diet_plans, consultations, food_items, health_logs, recipes, and forum_posts. Includes Row Level Security policy definitions for each table.

Appendix D – Application Screenshots

Contains screenshots of the deployed prototype including: Prakriti assessment screens, AI-generated diet plan views, food scanner interface, consultation booking flow, health progress tracker dashboards, AyurCare support chatbot, and the Ayur Forum community hub.

Appendix E – AWS CloudWatch and Grafana Dashboard Screenshots

Contains screenshots of the AWS CloudWatch monitoring dashboards and Grafana visualizations showing ECS task metrics, ALB request distributions, RDS performance, Lambda invocation counts, and cost analysis charts.